

Multi-Horizon Failure-Risk Models for Battery Failure Prediction: How Much Cross-Chemistry Transfer Is a State-of-Health Shortcut?

Anonymous

Affiliation 1, City, Country (e-mail: author1@example.com)

Abstract

Predicting whether a lithium-ion cell will fail inside a short operating window is a different question from estimating its remaining useful life, and it is the one that matters most for real-time dispatch in electric vehicles and grid storage. This study widens a recently proposed multi-horizon failure-risk classification framework, first shown on one model and one dataset, into a validity-focused evaluation covering four model families, a true discrete-time hazard model, trivial health-state baselines, and six transfer targets spanning four chemistries — LCO, LFP, nickel-rich NCA and NMC — drawn from six laboratories, all under cell-disjoint splits and a leakage-clean, cross-fitted calibration protocol. Within-dataset discrimination is reliable, with fold-mean area under the curve (AUC) of 0.90–0.96 across the tree ensembles under Platt calibration.

The central result is a controlled State-of-Health (SOH) ablation: apparent cross-dataset transfer with SOH included (pooled AUC up to 1.00) collapses without SOH (at best 0.58 at $H=20$), and much of the with-SOH signal is reproducible by a distance-to-threshold rule on SOH alone (0.912 AUC on the 141-cell MIT–Stanford Severson set). Because SOH also enters the failure definition and the label window includes the current cycle, all with-SOH numbers are diagnostic-proximity upper bounds rather than pure prognosis: relabeling failures by voltage sag alone, a criterion SOH does not define, cuts with-SOH transfer by 0.17 (0.90 to 0.72 pooled on Severson at $H=20$), and restricting evaluation to rows scored while SOH is still above 0.80 removes the concurrent-detection component everywhere. Same-chemistry cross-laboratory controls confirm that the collapse is driven by dataset shift at least as much as by chemistry, so SOH’s dominance is

a dataset- and chemistry-dependent health-state shortcut. The evaluation is replicated on 60 new cells of NMC, NCA and LFP chemistry from two further laboratories (the public Battery Archive): with-SOH transfer spans 0.77–1.00 pooled AUC and ranges down to 0.36–0.91 without SOH outside the single-condition HNEI group (paired deltas 0.06–0.64), and no competing shortcut hides in temperature or C-rate. Isotonic calibration largely fails to transfer, and the three tree ensembles deploy on a \$12 ESP32-S3 in under a millisecond per inference for a single SRAM-resident model (0.76–2.36 ms from flash for the full binary).

Keywords: Battery failure prediction, multi-horizon failure-risk classification, state-of-health shortcut, cross-chemistry transfer, dataset shift, probability calibration, discrete-time hazard model, SHAP, embedded machine learning, ESP32-S3, battery archive.

I. Introduction

Lithium-ion batteries underpin electrified transport and grids, yet field reliability remains hard to certify. Conventional prognostics casts the problem as remaining-useful-life (RUL) regression [1, 2], which answers a question operators rarely ask in real time: a battery management system needs to know whether a cell survives the next mission, not its aggregate cycle count. Predicting lifetime from field rather than laboratory data remains open [3]. Tree ensembles — Random Forest [4], XGBoost [5], LightGBM [6] — serve capacity and RUL estimation, but binary failure-risk classification at multiple horizons is recent [7].

Shikdar and Laaksonen [7] recast the problem as multi-horizon failure-risk classification: for cycle t and horizon H , predict whether failure (SOH below 0.80 or voltage sag beyond 6%) occurs within H cycles. Their gradient boosting model on NASA 18650 (37 LCO cells) reached AUC 0.87–0.90 across $H \in \{10, 20, 30, 50\}$, but left three gaps: one model on one dataset (no learner, hyperparameter, or calibration sensitivity), no cross-dataset check, and no path to the microcontroller where a battery management system runs.

This paper begins from those gaps but is organized around a sharper validity question: *when a failure-risk model appears to transfer between datasets and chemistries, is anything about degradation being transferred, or is the model merely reading a health-state variable that encodes proximity to the failure criterion?* The composite label is defined from SOH, and SOH is also

Table 1: Comparison of Representative Battery Studies With This Work.

Study	Family	Data	Position relative to this work
Meng and Li [1]	PHM review	Various	Physics-based, data-driven, hybrid families
Breiman, Chen, Ke [4, 5, 6]	Tree ensembles	RUL tasks	RUL regression, not failure-risk classification
Sahoo et al. [8]	Transfer (SOH)	NASA, CALCE	SOH regression only
Lu et al. [9]	Deep SOH	Large corpora	SOH regression only
Platt, Zadrozny et al. [10, 11, 12, 13]	Calibration	Various	In-domain comparison only
Lundberg et al. [15]	Explainability	—	TreeSHAP for ensembles
Grzesik and Mrozek [16]	Edge ML	—	No battery failure cases
Gupta and Ramdas [17]	Online calib.	—	Distribution-shift Platt
Shikdar and Laaksonen [7]	Failure risk	NASA	One model, one dataset (baseline)
This work	Failure risk + validity	8 datasets	Multi-model SOH-shortcut diagnosis, hazard model, edge deployment

the strongest feature, so a model can score well on transfer while learning only “distance to the SOH = 0.80 threshold.”

Two adjacent literatures bear on these questions. Transfer learning for battery state estimation is a growing area: Sahoo et al. [8] fine-tuned an SOH estimator across NASA and CALCE cells, and Lu et al. [9] pretrained on large-scale degradation data; both target SOH regression rather than failure-risk classification, where the horizon-dependent label and shifting calibration add complications. On calibration, Platt scaling [10] fits a sigmoid to raw scores while isotonic regression [11] fits a step function; Platt is safer on small datasets [12] and imbalanced data [13], but no prior study compares the two under cross-chemistry shift. A companion study found the calibration layer is the primary point of failure under operating perturbation and that small-sample recalibration largely restores it [14]; whether the repair survives cross-chemistry shift is examined here. TreeSHAP [15] is the standard attribution tool for ensembles, edge inference removes latency and connectivity needs [16], and online recalibration under shift is an active area [17]. Table 1 situates this work against these lines of research.

The objective of this study is to close the three gaps identified above and, in doing so, to answer the validity question rigorously. The contributions are:

- a four-model benchmark (three tree ensembles plus a GRU) on two LCO datasets under cell-disjoint splits with cross-fitted, leakage-free calibration;
- a discrete-time hazard formulation whose survival product is monotone

in the horizon by construction, against independent fixed-horizon classifiers that violate monotonicity on 32–45% of rows (see Supplementary Material);

- a shortcut diagnosis of cross-dataset transfer: minimal baselines, a factorial feature ablation, a failure-definition ablation relabeling failures by SOH or voltage sag alone, and same-chemistry cross-laboratory controls separating dataset shift from chemistry shift;
- a cell-as-unit uncertainty treatment: cell-level bootstrap intervals on every headline number, DeLong tests demoted to secondary evidence;
- calibration analysis under shift with small-sample recalibration, a cost-sensitive operational evaluation, and a complete ESP32-S3 embedded deployment.

The remainder of this paper is organized as follows. Section II describes the methodology, Section III the results, Section IV the analysis and limitations, and Section V concludes.

II. Methodology

II-A. Composite Failure Label

Following the original protocol [7], cell State-of-Health at cycle k is discharge capacity over mean first-10-cycle capacity,

$$\begin{aligned}\text{SOH}_k^{(c)} &= \frac{Q_k^{(c)}}{\bar{Q}^{(c)}}, \\ \bar{Q}^{(c)} &= \frac{1}{10} \sum_{i=1}^{10} Q_i^{(c)},\end{aligned}\tag{1}$$

where $Q_k^{(c)}$ is the discharge capacity at cycle k . The voltage baseline $\bar{V}^{(c)}$ is defined analogously. A discharge cycle t in cell c is labeled as failing within horizon H if any cycle $k \in [t, t+H)$ meets either of two criteria,

$$y_t^{(c,H)} = \begin{cases} 1, & \exists k \in [t, t+H) : \text{SOH}_k^{(c)} \leq 0.80 \text{ or} \\ & \bar{V}_k^{(c)} < 0.94 \bar{V}^{(c)}, \\ 0, & \text{otherwise,} \end{cases}\tag{2}$$

and once triggered the label stays 1 for all later cycles. Because the window $[t, t+H)$ includes the current cycle, rows scored at or after the failure crossing are labeled positive for every formulation in this paper: with-SOH results therefore mix concurrent failure detection with future failure prediction, and Section III-C reports the two components separately. This paragraph makes explicit what was already true of every evaluation reported here; the label code is unchanged and no reported number moves as a result of the clarification. The task is to estimate the cumulative failure risk over the horizon,

$$p_t^{(c,H)} = \Pr(T_f \leq t+H \mid T_f > t, \mathbf{x}_t^{(c)}), \quad (3)$$

where T_f is the failure cycle and $\mathbf{x}_t^{(c)}$ the feature vector at cycle t . Datasets without voltage data, or with voltage at the cutoff as in Oxford’s flat LFP plateau, fall back to SOH alone.

Three properties of this label definition matter for everything that follows. First, (3) is a fixed-horizon cumulative failure risk, not a discrete-time hazard: each horizon is an independent binary classifier with no monotonicity constraint $p_t^{(c,10)} \leq p_t^{(c,20)} \leq p_t^{(c,30)} \leq p_t^{(c,50)}$. Section II-G adds a true discrete-time hazard model, and Section III-H quantifies the violations. Second, the label is *partially defined by SOH itself*, and SOH is also an input feature; a with-SOH model can therefore detect proximity to the SOH = 0.80 threshold without learning degradation dynamics. Every with-SOH result here is an upper bound including this leakage, and Section III-F measures it by relabeling failures one endpoint at a time. Third, the endpoints separate: the label can use the SOH criterion alone, voltage sag alone, or both, independently of the feature ablation.

II-B. Datasets

Eight public aging datasets are used:

- **1) NASA 18650** [18]: 37 LCO cells aged under randomized charge–discharge profiles at room temperature, 1 028 cleaned discharge rows in total (about 28 per cell), with diverse failure modes;
- **2) CALCE LCO/CX2** [19]: 7 cells (CS2_33–36, CX2_36–38) at 1C/1C; ambient temperature and discharge duration unlogged;
- **3) Oxford** [20]: 5 LFP pouch cells at 1C/1C; discharge summaries at roughly 100-cycle intervals leave 46–78 usable cycles per cell (320 rows);

- 4) **MIT–Stanford Severson** [21, 22]: 141 LFP cells under a fast-charging protocol, 101–2 237 cleaned cycles each (about 117 000 rows in total);
- 5) **HNEI NMC/LCO** [23]: 15 NMC/LCO-blend 18650 cells cycled continuously at 25 °C under a fixed charge protocol and mixed discharge C-rates, 1016–1064 cycles each — all reaching the 80% SOH end-of-life criterion;
- 6) **SNL NCA** [23, 24]: 18 NCA cells at 15–35 °C, 451–904 cycles each;
- 7) **SNL NMC** [23, 24]: 6 NMC 18650 cells at 15–25 °C, 384–670 cycles each;
- 8) **SNL LFP** [23, 24]: 21 LFP cells at 15–35 °C, 3022–4544 cycles each (5 reach 80% SOH).

Items 5–8 are the Battery Archive “Cycle Test Data” export [23] (September 2026 summaries; SNL matrix per Preger et al. [24]), used only as *transfer targets* and within-group controls.

A pre-registered cleaning rule admits a cell only if it passes formation-cycle removal, Hampel despiking of the capacity trace, truncation at sustained level shifts, per-cycle voltage validity, and a minimum of 20 usable cycles. Of 75 downloaded cells, 60 pass (15 + 6 NMC, 18 NCA, 21 LFP; 106 787 cycle rows); the remaining 15 are excluded — all 15 for mixed state-of-charge windows — and are never used. On every admitted cell the SOH end-of-life crossing occurs no later than the voltage-sag crossing, so the composite label on the new groups is, in effect, the SOH criterion alone.

One consequence of the Oxford logging resolution deserves emphasis. Consecutive cleaned cycles are roughly 100 cycles apart, so every horizon $H \in \{10, 20, 30, 50\}$ is shorter than the sampling interval and the positive-label rows are *identical for all four horizons*: Oxford measures pooled failure detection, not horizon-resolved prediction, and its per-cell means rest effectively on four cells (Cell5 contributes one positive row). Both caveats apply to every Oxford number in this paper.

II-C. Features and Preprocessing

Each cycle is described by seven scalar features: cycle index, average and minimum discharge voltage, average discharge current, average temperature,

discharge duration, and SOH. Trees are scale-invariant and need no standardization; the GRU and logistic baselines are standardized with training-set statistics. Unlogged CALCE and Battery Archive columns (temperature, duration, current) are zero-filled; a constant-zero column cannot split within one dataset but acts as a dataset identifier in pooled training, a confound Section III-G controls with a common-feature subset.

II-D. Tree Models and Hyperparameters

Three tree ensembles are benchmarked with hyperparameters matched to the original study [7]: XGBoost [5] with `max_depth` 4, 300 estimators, learning rate 0.05, `subsample` 0.8, `colsample_bytree` 0.8, and `min_child_weight` 5; LightGBM [6] with the same depth, count, and rate; and Random Forest [4] with `max_depth` 6 and 300 estimators. All use `random_state` 42 for deterministic replication.

II-E. GRU Sequence Classifier

A compact GRU tests whether sequence-aware representations transfer across chemistries. It ingests sliding windows of $W=10$ cycles per cell through one 8-unit GRU layer,

$$\begin{aligned}\mathbf{r}_t &= \sigma(W_r \mathbf{x}_t + U_r \mathbf{h}_{t-1}), \\ \mathbf{z}_t &= \sigma(W_z \mathbf{x}_t + U_z \mathbf{h}_{t-1}), \\ \tilde{\mathbf{h}}_t &= \tanh(W_h \mathbf{x}_t + U_h(\mathbf{r}_t \odot \mathbf{h}_{t-1})), \\ \mathbf{h}_t &= (1 - \mathbf{z}_t) \odot \mathbf{h}_{t-1} + \mathbf{z}_t \odot \tilde{\mathbf{h}}_t,\end{aligned}\tag{4}$$

and the final hidden state maps through a sigmoid linear layer to the failure probability (Adam, learning rate 0.005, binary cross-entropy). The 8-unit width matches the small training-cell counts. For cross-dataset runs the feature set is reduced to the common, unit-safe set — cycle number, average and minimum voltage, plus SOH when enabled — standardized with training-set statistics; current, temperature, and duration are dropped for unit safety. The trees are evaluated on the identical common set (Section III-C), so no architecture claim rests on a feature-confounded comparison.

II-F. Minimal Baselines

Because the central question is whether apparent performance is a health-state shortcut, the most important models are embarrassingly simple ones [7]:

- **SOH distance rule:** the fitted-free score $d_t = 0.80 - \text{SOH}_t$, i.e., distance below the label threshold;
- **SOH-only** logistic regression, $\Pr(y_t^{(H)}=1) = \sigma(a_H \text{SOH}_t + b_H)$;
- **cycle-only** logistic regression on the cycle index;
- **SOH + cycle** logistic regression;
- **sensors-only** logistic regression on the five non-SOH, non-cycle features;
- **full-feature** logistic regression on all seven features.

If the SOH-only and distance-rule baselines reproduce most of the transferable performance of the tuned ensembles, then the ensembles contribute little beyond reading the health state — the negative result this validity analysis predicts.

II-G. A True Discrete-Time Hazard Model

To complement the four independent fixed-horizon classifiers, we also fit an explicit discrete-time survival model. For each observed cycle t and offset $\ell \in \{1, \dots, L\}$ with $L=50$, define the hazard

$$h_\ell(t) = \Pr(T_f = t+\ell \mid T_f \geq t+\ell, \mathbf{x}_t), \quad (5)$$

conditional on the features measured at t only. One pooled classifier — XGBoost, and a logistic regression — is fit on the offset-expanded at-risk rows $(\mathbf{x}_t, \ell)$ with target $\mathbf{1}\{T_f = t+\ell\}$, restricted to $t+\ell \leq T_f$, with ℓ as an input feature. The cumulative risk over horizon H is the survival product

$$R(t, H) = 1 - \prod_{\ell=1}^H (1 - \hat{h}_\ell(t)), \quad (6)$$

which is monotone non-decreasing in H by construction and yields all four horizons from one model. Rows at or after failure are scored exactly like the fixed-horizon classifiers score them, so the evaluation rows are identical across formulations.

II-H. Calibration Protocol

Three post-hoc calibration schemes are compared: Platt scaling fits a sigmoid to the raw scores s_i ,

$$p_i = \frac{1}{1 + \exp(-\alpha s_i - \beta)}, \quad (7)$$

by logistic regression; isotonic regression [11] fits the non-decreasing step function that minimizes squared error, implemented with scikit-learn’s `IsotonicRegression` and out-of-bounds clipping; and temperature scaling divides the logit of the score by a single fitted temperature.

The protocol is leakage-clean. For every within-dataset fold, calibrators are fitted on *cross-fitted out-of-fold scores* from an inner cell-disjoint Group-KFold over the training cells; untouched test cells are scored by a model trained on all training cells. In transfer, calibrators are fitted on source out-of-fold scores only. No calibrator is ever fitted on the scores it is applied to, and no test cell contributes to training or calibration.

A reproducibility note is retained from the original study comparison. The original study [7] reported Brier scores around 0.032; re-execution yields roughly 0.09–0.29, a gap of unknown cause that does not affect the relative comparisons forming the core of this paper.

II-I. Transfer Protocols

Transfer is evaluated by training on LCO cells (NASA alone, CALCE alone, or both) and testing on the target cells. Performance is reported per cell and pooled with cell-level bootstrap intervals; pooling alone can hide degenerate constant-score behavior, which per-cell statistics expose. Three families of transfer are used:

- **Cross-chemistry** (LCO→LFP): NASA, CALCE, or both → Oxford and → Severson, with and without SOH, and on the common feature set shared with the GRU;
- **Same-chemistry, cross-laboratory controls**: NASA→CALCE and CALCE→NASA (LCO→LCO) and Severson→Oxford and Oxford→Severson (LFP→LFP). Chemistry, laboratory, protocol, and logging all change at once between any source and target, so “cross-chemistry” is shorthand for cross-*dataset* shifts, and these controls separate the two;
- **Feature-definition ablations**: the factorial feature grid of Section III-E and the failure-endpoint relabeling of Section III-F.

II-J. Evaluation and Statistical Testing

Within-dataset evaluation uses 5-fold cell-grouped cross-validation on NASA (37 cells) and leave-one-cell-out on CALCE (7 cells): splits are cell-disjoint, models retrained per horizon. The experimental unit is the cell: headline AUCs are fold means with fold SDs, every headline carries a 95% cell-resampling interval (2 000 draws, 400 for Severson), and paired comparisons reuse the same resamples.

The DeLong nonparametric test [25] for paired ROC curves is retained as secondary evidence. It treats every cycle-level prediction as independent; because cycles within a cell are correlated, its p -values are inflated by pseudoreplication and are read only as directional consistency checks, never as precise significance statements.

Discrimination alone is not enough for a reliability-oriented study, so each within-dataset configuration also reports the Brier score, expected calibration error (10 equal-width bins), log loss, the slope and intercept of a logistic recalibration of the predicted probabilities (slope 1 and intercept 0 indicate perfect calibration), the area under the precision–recall curve, which is sensitive to the horizon-dependent class prevalence that AUC ignores, and sensitivity at a fixed 10% false-positive rate. Operational, cost-sensitive evaluation is described with the results in Section III-K.

Three reporting conventions apply throughout. Pooled AUC is dominated by long-lived cells, so per-cell statistics and cell-level bootstrap intervals are the primary evidence. Failure prevalence rises with the horizon, so Brier score, calibration error, thresholds, and precision–recall are horizon-dependent. The cost model prices a missed failure at 20 times a false alarm; thresholds use a 10% source false-alarm budget, and sensitivity to other budgets is untested.

II-K. Embedded Deployment Architecture

The deployment pipeline exports the three trained models into one flat binary run on an ESP32-S3 by a hand-written C tree walker: a 12-byte header, a 3×4-byte offset table, and per-model sections of consecutive node arrays (split feature, float threshold, child indices; leaves store class-1 probabilities). Writing ℓ_t for the leaf value reached in tree t and T for the number of trees, Random Forest averages and clamps its leaf probabilities,

$$p = \text{clamp}\left(\frac{1}{T} \sum_{t=1}^T \ell_t, 0, 1\right), \quad (8)$$

Table 2: Within-Dataset Failure-Risk Performance Across Four Horizons (Fold-Mean AUC $\pm$ Fold SD; Tree Ensembles and GRU Platt-Calibrated, Hazard Models Raw; CALCE Uses Leave-One-Cell-Out, NASA 5-Fold Cell-Grouped CV). Brier Is the Platt Pooled Value Averaged Over Horizons (Raw for Hazard Rows).

Model	Dataset	H=10	H=20	H=30	H=50	Mean AUC	Brier
XGBoost	NASA	0.903 $\pm$ 0.097	0.899 $\pm$ 0.121	0.910 $\pm$ 0.111	0.918 $\pm$ 0.115	0.907	0.182
LightGBM	NASA	0.912 $\pm$ 0.065	0.909 $\pm$ 0.126	0.892 $\pm$ 0.121	0.901 $\pm$ 0.073	0.903	0.206
Random Forest	NASA	0.911 $\pm$ 0.106	0.901 $\pm$ 0.109	0.912 $\pm$ 0.113	0.927 $\pm$ 0.107	0.913	0.201
GRU	NASA	0.872 $\pm$ 0.122	0.893 $\pm$ 0.114	0.884 $\pm$ 0.083	0.856 $\pm$ 0.128	0.876	0.307
Hazard XGBoost	NASA	0.918 $\pm$ 0.086	0.935 $\pm$ 0.086	0.937 $\pm$ 0.081	0.927 $\pm$ 0.070	0.929	0.182
Hazard Logistic	NASA	0.902 $\pm$ 0.064	0.899 $\pm$ 0.057	0.896 $\pm$ 0.045	0.867 $\pm$ 0.052	0.891	0.175
XGBoost	CALCE	0.965 $\pm$ 0.041	0.967 $\pm$ 0.036	0.967 $\pm$ 0.040	0.943 $\pm$ 0.071	0.961	0.075
LightGBM	CALCE	0.967 $\pm$ 0.047	0.968 $\pm$ 0.038	0.943 $\pm$ 0.070	0.942 $\pm$ 0.076	0.955	0.075
Random Forest	CALCE	0.934 $\pm$ 0.074	0.933 $\pm$ 0.070	0.917 $\pm$ 0.098	0.903 $\pm$ 0.110	0.922	0.070
GRU	CALCE	0.952 $\pm$ 0.072	0.959 $\pm$ 0.083	0.935 $\pm$ 0.090	0.921 $\pm$ 0.090	0.942	0.130
Hazard XGBoost	CALCE	0.894 $\pm$ 0.135	0.893 $\pm$ 0.130	0.916 $\pm$ 0.095	0.889 $\pm$ 0.096	0.898	0.108
Hazard Logistic	CALCE	0.979 $\pm$ 0.042	0.948 $\pm$ 0.119	0.936 $\pm$ 0.138	0.945 $\pm$ 0.117	0.952	0.122

while XGBoost and LightGBM pass the summed score through a sigmoid,

$$p = \frac{1}{1 + \exp(-\sum_{t=1}^T \ell_t - s_0)}, \quad (9)$$

where s_0 is the model’s initialization score ($s_0 \neq 0$ only for XGBoost). Correctness is confirmed in three stages — a Python walker mirroring the C logic against `predict_proba()`, a cross-compiled x86_64 binary, and on-target ESP32-S3 execution — all below 10^{-5} maximum absolute error before deployment is accepted.

III. Results

III-A. Within-Dataset Performance

Table 2 summarizes within-dataset performance as fold-mean AUC $\pm$ SD under the leakage-clean protocol (LOCO on CALCE, 5-fold grouped CV on NASA), with Platt Brier scores. Every tree configuration clears 0.89 at every horizon on both datasets (Fig. 1); the GRU matches the means with higher fold variance, and the single-fit hazard model lands in the same range. On the four new Battery Archive groups, 38 of 39 tree-group-horizon configurations reach 0.85, spanning 0.61–1.00 — though the mostly-censored LFP-SNL group (0.031 prevalence) is supporting evidence only.




Figure 1: Within-dataset AUC heatmap for the tree ensembles, mean across horizons, Platt-calibrated. Right block: the four new Battery Archive groups. GRU and hazard within-dataset numbers are in Table 2.

III-B. Calibration Under Cross-Fitting

Table 3 compares the three calibrators, averaged over the tree models and horizons, where every calibrator is fitted on cross-fitted out-of-fold scores (Section II-H). Three observations follow. First, on discrimination the raw scores are effectively unbeatable within dataset: any fitted calibrator is a monotone re-mapping and preserves per-fold ranking up to the ties it introduces, and the fold-mean AUC differences between calibrators are correspondingly small. Second, the step function of isotonic regression creates large tied groups whose cost is invisible to fold-mean ROC-AUC but severe for precision–recall: on NASA, raw scores reach a mean PR-AUC of 0.774 while isotonic collapses to 0.639 (Fig. 3). Third, calibration quality proper is best for Platt on CALCE (ECE 0.029 against 0.049 for isotonic), but the fitted Platt slope on CALCE is only 0.011: on that long-tailed, leave-one-cell-out setting the out-of-fold sigmoid essentially washes confidence out toward the base rate while preserving ranking. Temperature scaling is the most balanced compromise on NASA (slope 1.218, Brier 0.178, both better than Platt’s). There is no global winner across datasets: temperature leads on NASA, Platt on CALCE ECE at the cost of a destroyed slope. We therefore read the comparison as: *Platt preserves discrimination best while achieving calibration error comparable to or better than isotonic*, which is a materially weaker and more defensible claim than “Platt is the better calibrator,” and raw scores remain the safest ranking signal whenever the operating context is uncertain.

Table 3: Calibration Quality Averaged Over Tree Models and Horizons. All Calibrators Are Fitted on Cross-Fitted Out-of-Fold Scores. Slope 1 / Intercept 0 Indicate Perfect Probabilistic Calibration.

Dataset	Method	AUC (fold mean)	Brier	ECE	Log loss	Slope	Intercept
NASA	Isotonic	0.854	0.207	0.151	0.652	0.48	-0.26
NASA	Platt	0.908	0.196	0.136	0.575	0.90	-0.14
NASA	Temperature	0.908	0.178	0.107	0.536	1.22	-0.30
CALCE	Isotonic	0.915	0.067	0.049	0.279	0.41	1.17
CALCE	Platt	0.946	0.073	0.029	0.280	0.01	2.54
CALCE	Temperature	0.945	0.080	0.067	0.284	0.53	1.42




Figure 2: Reliability diagrams under the cross-fitted protocol. (a) Within CALCE, XGBoost at $H=20$: the fitted Platt sigmoid preserves ranking but washes confidence toward the base rate. (b) ALL-LCO → Severson without SOH at $H=20$: no fitted calibrator, whichever source scores it is applied to, repairs calibration on the shifted target.

III-C. Cross-Dataset Transfer With and Without SOH

Table 4 reports LCO→LFP transfer at $H=20$ with SOH kept, and Table 5 without; Fig. 4 maps the AUC drop from removing SOH. Every number carries a cell-level bootstrap interval plus per-cell mean and standard deviation. With SOH, the trees reach pooled raw AUC up to 1.00 (Oxford) and 0.90 (Severson). Without SOH, transfer collapses to at most 0.58 on Oxford and 0.53–0.75 on Severson — point estimates fall below 0.5 in several configurations, significantly so only for the hazard model on Oxford.

The GRU shows a source-dependent instability at three seeds: per-cell AUC on Severson without SOH is 0.987 ± 0.015 for NASA (stable, near-oracle) but 0.479 ± 0.451 for CALCE (seeds 0.21–1.00) and 0.893 ± 0.131 pooled — against 0.998 ± 0.003 , 0.964 ± 0.033 and 0.991 ± 0.015 with SOH.




Figure 3: Within-dataset PR-AUC of raw scores versus horizon. Precision–recall performance, unlike ROC-AUC, is sensitive to the horizon-dependent failure prevalence; it is also where isotonic’s tied step outputs pay for their ranking-free compression.

Table 4: Cross-Dataset Transfer at $H=20$ With SOH (Upper-Bound / Label-Proximity Regime). Raw-Score Pooled AUC With Cell-Level Bootstrap 95% CI; Per-Cell Mean $\pm$ SD Across Test Cells (Oxford: Five Cells, One With a Single Positive Row; See Section II-B).

Training	Model	Oxford		Severson	
		Pooled AUC [CI]	Per-cell	Pooled AUC [CI]	Per-cell
NASA	XGBoost	0.957 [0.931, 0.987]	0.959 $\pm$ 0.037	0.871 [0.767, 0.997]	0.997 $\pm$ 0.010
NASA	LightGBM	1.000 [1.000, 1.000]	1.000 $\pm$ 0.000	0.849 [0.724, 0.998]	0.992 $\pm$ 0.048
NASA	Random Forest	1.000 [1.000, 1.000]	1.000 $\pm$ 0.000	0.861 [0.739, 0.998]	0.999 $\pm$ 0.007
CALCE	XGBoost	0.850 [0.764, 0.910]	0.763 $\pm$ 0.245	0.850 [0.749, 0.967]	0.957 $\pm$ 0.054
CALCE	LightGBM	0.889 [0.827, 0.932]	0.837 $\pm$ 0.193	0.863 [0.782, 0.957]	0.952 $\pm$ 0.037
CALCE	Random Forest	0.850 [0.773, 0.899]	0.766 $\pm$ 0.169	0.902 [0.834, 0.984]	0.996 $\pm$ 0.022
ALL LCO	XGBoost	0.917 [0.878, 0.965]	0.826 $\pm$ 0.265	0.896 [0.812, 0.995]	0.995 $\pm$ 0.015
ALL LCO	LightGBM	0.971 [0.959, 0.981]	0.969 $\pm$ 0.015	0.882 [0.789, 0.990]	0.993 $\pm$ 0.026
ALL LCO	Random Forest	0.989 [0.979, 0.998]	0.992 $\pm$ 0.013	0.889 [0.796, 0.998]	0.996 $\pm$ 0.021

On the identical common feature set the trees retain their table 4 behavior, so the gap is architectural, not featural. On Oxford the GRU is at or below chance (Table 6).

The reading that survives the uncertainty treatment is deliberately modest. SOH appears on both sides of the learning problem — in the feature vector and inside the label definition — so high with-SOH transfer is com-

Table 5: Cross-Dataset Transfer at H=20 Without SOH (Prognostically Meaningful Regime). Columns as in Table 4.

Training	Model	Oxford		Severson	
		Pooled AUC [CI]	Per-cell	Pooled AUC [CI]	Per-cell
NASA	XGBoost	0.518 [0.513, 0.522]	0.519±0.006	0.595 [0.510, 0.726]	0.676±0.184
NASA	LightGBM	0.533 [0.497, 0.543]	0.448±0.201	0.527 [0.422, 0.672]	0.535±0.183
NASA	Random Forest	0.521 [0.487, 0.537]	0.442±0.198	0.674 [0.561, 0.826]	0.836±0.176
CALCE	XGBoost	0.416 [0.264, 0.554]	0.360±0.208	0.682 [0.670, 0.695]	0.806±0.082
CALCE	LightGBM	0.462 [0.307, 0.579]	0.410±0.207	0.643 [0.530, 0.721]	0.830±0.042
CALCE	Random Forest	0.501 [0.320, 0.641]	0.435±0.243	0.665 [0.547, 0.749]	0.870±0.049
ALL LCO	XGBoost	0.429 [0.279, 0.597]	0.415±0.240	0.750 [0.692, 0.835]	0.860±0.139
ALL LCO	LightGBM	0.332 [0.112, 0.515]	0.316±0.209	0.718 [0.634, 0.850]	0.808±0.167
ALL LCO	Random Forest	0.581 [0.516, 0.644]	0.608±0.103	0.746 [0.673, 0.845]	0.876±0.115




Figure 4: Transferability collapse map at $H=20$ for the tree ensembles: the AUC drop from removing SOH, per source→target pair and model. Positive values mean SOH removal destroys apparent transfer performance; the near-uniform saturation at the maximum drop is the signature of a single-feature shortcut rather than of distributed degradation knowledge. Right block: ALL-LCO → new Battery Archive targets. GRU deltas are reported in the text.

Table 6: GRU Transfer at H=20 on the Common Feature Set: Seed-Averaged Per-Cell AUC ($\pm$ SD Over Seeds 42, 1, 7).

Training	Feature set	Oxford	Severson
NASA	common w/ SOH	0.196±0.155	0.998±0.003
NASA	common no SOH	0.147±0.210	0.987±0.015
CALCE	common w/ SOH	0.182±0.311	0.964±0.033
CALCE	common no SOH	0.099±0.171	0.479±0.451
ALL LCO	common w/ SOH	0.088±0.084	0.991±0.015
ALL LCO	common no SOH	0.000±0.000	0.893±0.131

patible with two mechanisms: a chemistry-dependent mapping from SOH to remaining life, and mere detection of proximity to the $\text{SOH} = 0.80$ criterion that defines the label. The experiments that follow separate these mechanisms.

Diagnosis or prognosis? Because the label window includes the scoring cycle, rows from already-failed cells score positive. Table 7 splits headline transfer numbers into all-rows AUC and prognosis-only AUC on the same scores restricted to rows with SOH still above 0.80. On Severson at $H=20$, XGBoost with SOH moves from 0.896 to 0.838 [0.760, 0.991] and without SOH from 0.750 to 0.721 [0.666, 0.833]: 99% of rows are scored above threshold, so concurrent detection explains only about 0.05–0.07 there; the rest is near-threshold ranking, priced by the failure-definition ablation in Section III-F. The distance rule behaves identically (0.912 to 0.862). On the new targets the diagnostic share is larger wherever cells linger below threshold — HNEI 0.993 to 0.963 (28

III-D. Minimal Baselines: How Much Is the SOH Shortcut?

Table 8 places the tuned ensembles next to trivial baselines. The distance-to-threshold rule $d_t = 0.80 - \text{SOH}_t$, which involves no fitting at all, transfers to Severson at pooled AUC 0.912 [CI] for the ALL-LCO source — matching or exceeding the tuned ensembles’ with-SOH transfer in Table 4. An SOH-only logistic regression behaves equivalently (0.912). Both are diagnostic baselines by construction: they order cells by proximity to the label threshold, not by learned degradation dynamics. By contrast, the sensor-only logistic regression transfers at 0.301 on Severson (NASA source): below chance, meaning the remaining sensor features actively invert on the shifted distribution. Within CALCE, the cycle index alone ranks failure at pooled AUC 0.936, approaching the tuned ensembles, because CALCE’s seven cells span very different cycle lives; on NASA, whose cells age under randomized profiles, the same feature is inverted (0.354).

Three conclusions follow. First, most of what transfers is the SOH coordinate, not degradation knowledge. Second, “more features” does not help transfer; sensor features carry dataset-specific artifacts that can invert. Third, the ensembles’ within-dataset value is undiminished — but the transfer claim must be downgraded from “models transfer” to “the health state transfers, and models can read it.”

Table 7: Diagnosis Versus Prognosis at H=20 (ALL-LCO Source, Raw Scores): All-Rows Pooled AUC [Cell-Level CI Where Computed] Against Prognosis-Only AUC on Rows Scored While SOH>0.80, With the Share of Rows Scored Above Threshold. GRU and Distance-Rule Rows Are Point Estimates.

Target	Model	Features	All rows	Prognosis-only	Frac. above
Severson	GRU	common, no SOH	0.631	0.553	0.99
Severson	GRU	common+SOH	0.827	0.756	0.99
Severson	LightGBM	no SOH	0.718 [0.634,0.846]	0.674 [0.597,0.851]	0.99
Severson	LightGBM	with SOH	0.882 [0.786,0.990]	0.814 [0.732,0.978]	0.99
Severson	SOH-dist rule	dist rule	0.912	0.862	0.99
Severson	SOH-dist rule	SOH-only	0.912	0.862	0.99
Severson	Random Forest	no SOH	0.746 [0.673,0.845]	0.705 [0.632,0.843]	0.99
Severson	Random Forest	with SOH	0.889 [0.783,0.998]	0.825 [0.737,0.996]	0.99
Severson	XGBoost	no SOH	0.750 [0.692,0.835]	0.721 [0.666,0.833]	0.99
Severson	XGBoost	with SOH	0.896 [0.804,0.995]	0.838 [0.760,0.991]	0.99
HNEI NMC	LightGBM	no SOH	0.988 [0.974,0.996]	0.905 [0.863,0.944]	0.28
HNEI NMC	LightGBM	with SOH	0.999 [0.998,0.999]	0.969 [0.946,0.987]	0.28
HNEI NMC	Random Forest	no SOH	0.996 [0.994,0.999]	0.945 [0.908,0.978]	0.28
HNEI NMC	Random Forest	with SOH	0.999 [0.999,1.000]	0.981 [0.970,0.993]	0.28
HNEI NMC	XGBoost	no SOH	0.989 [0.977,0.997]	0.908 [0.870,0.945]	0.28
HNEI NMC	XGBoost	with SOH	0.993 [0.983,0.999]	0.963 [0.935,0.987]	0.28
SNL NCA	LightGBM	no SOH	0.850 [0.791,0.911]	0.841 [0.798,0.889]	0.63
SNL NCA	LightGBM	with SOH	0.985 [0.974,0.993]	0.891 [0.845,0.946]	0.63
SNL NCA	Random Forest	no SOH	0.911 [0.875,0.951]	0.837 [0.779,0.907]	0.63
SNL NCA	Random Forest	with SOH	0.983 [0.974,0.992]	0.901 [0.858,0.957]	0.63
SNL NCA	XGBoost	no SOH	0.878 [0.824,0.930]	0.849 [0.803,0.895]	0.63
SNL NCA	XGBoost	with SOH	0.972 [0.956,0.988]	0.886 [0.837,0.943]	0.63
SNL NMC	LightGBM	no SOH	0.693 [0.570,0.959]	0.728 [0.577,0.956]	0.60
SNL NMC	LightGBM	with SOH	0.847 [0.743,0.982]	0.722 [0.604,0.956]	0.60
SNL NMC	Random Forest	no SOH	0.742 [0.689,0.874]	0.649 [0.604,0.793]	0.60
SNL NMC	Random Forest	with SOH	0.858 [0.748,0.975]	0.706 [0.613,0.929]	0.60
SNL NMC	XGBoost	no SOH	0.709 [0.612,0.953]	0.730 [0.599,0.946]	0.60
SNL NMC	XGBoost	with SOH	0.772 [0.677,0.979]	0.729 [0.608,0.960]	0.60
SNL LFP	LightGBM	no SOH	0.359 [0.282,0.443]	0.383 [0.311,0.467]	0.97
SNL LFP	LightGBM	with SOH	0.999 [0.998,1.000]	0.993 [0.988,0.998]	0.97
SNL LFP	Random Forest	no SOH	0.666 [0.569,0.769]	0.588 [0.565,0.616]	0.97
SNL LFP	Random Forest	with SOH	0.999 [0.998,1.000]	0.994 [0.990,0.999]	0.97
SNL LFP	XGBoost	no SOH	0.496 [0.455,0.531]	0.517 [0.481,0.575]	0.97
SNL LFP	XGBoost	with SOH	0.998 [0.997,1.000]	0.991 [0.985,0.997]	0.97

III-E. Factorial Feature Ablation

Table 9 runs the full $2^3 - 1$ grid over {SOH, cycle index, sensors} at H=20. Three regularities stand out. SOH-only is the strongest single group on both LFP targets, confirming that the health coordinate dominates whatever transfer exists. Cycle-only is competitive on Oxford in the logistic base-lines (0.99 NASA-source) — whose five cells fail at well-separated raw cycle counts, so the cycle index is a between-cell aging proxy — but not in the facto-

Table 8: Minimal Transfer Baselines at H=20: Pooled AUC With Cell-Level Bootstrap 95% CI. The Distance Rule Is the Unfitted Score 0.80—SOH.

Baseline	Oxford			Severson		
	NASA src	CALCE src	ALL src	NASA src	CALCE src	ALL src
SOH distance rule	1.000 [1.000,1.000]	1.000 [1.000,1.000]	1.000 [1.000,1.000]	0.912 [0.836,0.998]	0.912 [0.836,0.998]	0.912 [0.836,0.998]
SOH only	1.000 [1.000,1.000]	1.000 [1.000,1.000]	1.000 [1.000,1.000]	0.912 [0.836,0.998]	0.912 [0.836,0.998]	0.912 [0.836,0.998]
cycle only	0.989 [0.983,1.000]	0.756 [0.732,0.780]	0.787 [0.760,0.814]	0.713 [0.591,0.798]	0.713 [0.591,0.798]	0.713 [0.591,0.798]
SOH + cycle	0.990 [0.985,1.000]	0.797 [0.770,0.825]	0.818 [0.788,0.848]	0.865 [0.834,0.902]	0.764 [0.674,0.829]	0.758 [0.664,0.825]
sensors only	0.500 [0.500,0.500]	0.500 [0.500,0.500]	0.500 [0.500,0.500]	0.301 [0.015,0.549]	0.498 [0.497,0.500]	0.303 [0.059,0.497]
full (7 feats)	0.500 [0.500,0.500]	0.500 [0.500,0.500]	0.500 [0.500,0.500]	0.823 [0.675,0.995]	0.712 [0.589,0.797]	0.674 [0.525,0.782]

Table 9: Factorial Feature Ablation at H=20 (XGBoost, Pooled AUC With Cell-Level Bootstrap 95% CI). Feature Groups: SOH; Cycle Index; Sensors = Voltages, Current, Temperature, Duration.

Feature set	Oxford			Severson		
	NASA src	CALCE src	ALL src	NASA src	CALCE src	ALL src
full	0.957 [0.931,0.987]	0.850 [0.764,0.910]	0.917 [0.878,0.965]	0.871 [0.767,0.997]	0.850 [0.749,0.967]	0.896 [0.812,0.995]
no SOH	0.518 [0.513,0.522]	0.416 [0.264,0.554]	0.429 [0.279,0.597]	0.595 [0.510,0.726]	0.682 [0.670,0.695]	0.750 [0.692,0.835]
—cycle	0.965 [0.942,0.989]	0.967 [0.929,0.981]	0.993 [0.954,1.000]	0.854 [0.733,0.998]	0.766 [0.584,0.974]	0.774 [0.578,0.998]
—SOH—cycle	0.658 [0.522,0.797]	0.433 [0.282,0.594]	0.388 [0.224,0.506]	0.531 [0.433,0.680]	0.500 [0.500,0.500]	0.535 [0.455,0.653]
SOH only	0.965 [0.942,0.989]	1.000 [1.000,1.000]	1.000 [1.000,1.000]	0.834 [0.695,0.991]	0.798 [0.638,0.978]	0.833 [0.689,0.993]
cycle only	0.541 [0.537,0.545]	0.541 [0.537,0.545]	0.541 [0.537,0.545]	0.709 [0.695,0.724]	0.694 [0.681,0.708]	0.719 [0.704,0.735]
sensors	0.658 [0.522,0.797]	0.433 [0.282,0.594]	0.388 [0.224,0.506]	0.531 [0.433,0.680]	0.500 [0.500,0.500]	0.535 [0.455,0.653]

rial XGBoost grid itself (0.54 Oxford, 0.72 Severson): the pattern is model- and source-specific, not portable degradation knowledge. No combination without SOH reaches 0.76 on either target (best 0.75), and removing both SOH and the cycle index leaves the sensors with essentially nothing (0.535 with the ALL-LCO source on Severson). The ablation makes the shortcut diagnosis feature-by-feature: it is not that the models fail to use the sensor features; it is that the sensor features themselves carry no cross-dataset signal.

III-F. Failure-Definition Ablation: Measuring the Label Circularity

Table 10 relabels failure by a single endpoint at a time and repeats the transfer experiments. This quantifies the circularity concern of Section II-A: with-SOH transfer *tracks how much SOH defines the label*. Under the SOH-only label, transfer to Severson is 0.985; under the combined label, 0.896. Relabel by voltage sag alone and with-SOH transfer falls to 0.723 — a drop of about 0.17, with the with-minus-without gap narrowing from +0.15 to +0.11. The residual (0.723 versus 0.611) is genuine but weak health-state information. Within one laboratory the distinction nearly vanishes — on NASA the with-SOH model reaches 0.936 — because sensor features co-move with SOH there; dataset shift destroys that co-movement, and with it

Table 10: Failure-Definition Ablation at H=20 (XGBoost, ALL-LCO Source, Pooled AUC [CI]): Transfer With and Without SOH as a Feature, for Three Label Endpoints.

Label endpoint	Oxford		Severson	
	with SOH	no SOH	with SOH	no SOH
Combined (SOH or sag)	0.917 [0.878,0.965]	0.429 [0.279,0.597]	0.896 [0.812,0.995]	0.750 [0.692,0.835]
SOH only	0.995 [0.988,1.000]	0.464 [0.320,0.596]	0.985 [0.977,0.992]	0.759 [0.709,0.801]
Voltage sag only	—	—	0.723 [0.705,0.740]	0.611 [0.585,0.637]

Table 11: Same-Chemistry Cross-Laboratory Controls at H=20 (XGBoost, Pooled AUC [CI]). Full Feature Sets Exercise the Zero-Imputation Confound; Common Sets Restrict Both Datasets to the Shared Four-Feature Subset.

Pair	with SOH (full)	no SOH (full)	no SOH (common)
NASA→CALCE (LCO→LCO)	0.870 [0.799,0.966]	0.849 [0.753,0.985]	0.781 [0.708,0.964]
CALCE→NASA (LCO→LCO)	0.543 [0.356,0.677]	0.566 [0.427,0.678]	0.601 [0.496,0.703]
Severson→Oxford (LFP→LFP)	0.997 [0.991,1.000]	0.471 [0.466,0.477]	0.428 [0.422,0.435]
Oxford→Severson (LFP→LFP)	0.688 [0.605,0.788]	0.500 [0.500,0.500]	0.500 [0.500,0.500]

the transfer illusion.

III-G. Same-Chemistry Cross-Laboratory Controls

If the collapse of Section III-C were a chemistry effect, transfer should survive when chemistry is held constant. Table 11 shows it does not. NASA→CALCE (both LCO) retains most of its transfer without SOH, so some laboratory pairs do share portable sensor signal; but Severson→Oxford (both LFP) reproduces the full collapse pattern of the cross-chemistry shifts, including high with-SOH transfer and near-chance sensor-only transfer. These controls tighten the central attribution: the failure to transfer is a property of *dataset shift* — laboratories, protocols, and logging conventions — that chemistry change adds to but does not uniquely cause. SOH is therefore a *dataset- and chemistry-dependent* health-state signal, not a chemistry-specific one, and chemistry is not isolated as a causal variable.

III-H. Discrete-Time Hazard Model Versus Fixed-Horizon Classifiers

Table 12 compares the survival-product model of Section II-G against the fixed-horizon XGBoost classifier. Within datasets the two match to within a few AUC points at every horizon (Table 2); no formal equivalence test is run. Under transfer the survival product inherits the same SOH dependence — model formulation does not exempt inputs from dataset shift — but $R(t, H)$

Table 12: Discrete-Time Hazard Model at $H=20$: Transfer Performance (Pooled AUC [CI]). The Fixed-Horizon XGBoost Comparator With SOH Is in Table 4.

Training	Model	Oxford	Severson
NASA	Hazard XGBoost	0.971 [0.953,0.991]	0.884 [0.786,0.998]
NASA	Hazard Logistic	0.500 [0.500,0.500]	0.607 [0.402,0.890]
CALCE	Hazard XGBoost	0.466 [0.301,0.578]	0.737 [0.711,0.774]
CALCE	Hazard Logistic	0.500 [0.500,0.500]	0.500 [0.500,0.500]
ALL LCO	Hazard XGBoost	0.314 [0.135,0.457]	0.866 [0.770,0.978]
ALL LCO	Hazard Logistic	0.500 [0.500,0.500]	0.916 [0.873,0.969]

is monotone in H by construction, whereas the four independent classifiers violate $p_{10} \leq p_{20} \leq p_{30} \leq p_{50}$ on 32–45% of within-dataset rows (raw scores; Table 13), rising to 95.8% (nasa-trained xgboost) on transferred Severson rows. Monotone cumulative risk is what makes “risk within 50 cycles” coherent to quote: the hazard model buys structural coherence and a single fit at no within-dataset cost, and pays the same SOH shortcut as every other model.

III-I. Failure of Calibration Transfer, and Its Partial Repair

A second, independent failure mode appears in the cross-dataset tables when calibrated outputs are required. Isotonic regression is fitted as a step function on the source distribution; when target scores shift, many target scores fall into one bin and receive identical calibrated probabilities, and pooled ranking degrades accordingly — in Table 4, isotonic calibration costs the tree ensembles up to 0.48 pooled AUC on Oxford and 0.31 on Severson even with SOH present. Platt’s smooth sigmoid degrades more gracefully, which is precisely the sense in which “Platt better preserves discrimination at comparable calibration error.” Fig. 2(b) shows the deeper problem: on the shifted target, *no* source-fitted calibrator lands near the diagonal, however it is fitted.

The recalibration arms of the companion study [14] ask whether a small labeled target sample repairs this. Arm A refits a calibrator on k labeled target cells while leaving the classifier’s scores untouched; Arm B continues training the classifier on them. On Severson without SOH, zero-shot ECE is 0.817; at $k=5$ labeled cells, isotonic and Platt reduce it to 0.044 and 0.045, while temperature scaling remains poorly calibrated (0.563). Calibra-

Table 13: Monotonicity Violations: Share of Scored Rows With $p_{10} \leq p_{20} \leq p_{30} \leq p_{50}$ Failing, on Raw Scores (Platt in Parentheses). The Survival Product Is Monotone by Construction (Rate 0).

Setting	Target	Model	Raw	Platt
within	calce	XGBoost	0.452	0.056
within	nasa	XGBoost	0.323	0.260
transfer (calce src)	oxford	LightGBM	0.641	0.004
transfer (calce src)	oxford	Random Forest	0.501	0.540
transfer (calce src)	oxford	XGBoost	0.565	0.003
transfer (nasa+calce src)	oxford	LightGBM	0.210	0.648
transfer (nasa+calce src)	oxford	Random Forest	0.335	0.305
transfer (nasa+calce src)	oxford	XGBoost	0.609	0.341
transfer (nasa src)	oxford	LightGBM	0.521	0.468
transfer (nasa src)	oxford	Random Forest	0.559	0.376
transfer (nasa src)	oxford	XGBoost	0.248	0.124
transfer (calce src)	severson	LightGBM	0.405	0.013
transfer (calce src)	severson	XGBoost	0.420	0.072
transfer (nasa src)	severson	LightGBM	0.582	0.537
transfer (nasa src)	severson	Random Forest	0.778	0.557
transfer (nasa src)	severson	XGBoost	0.958	0.923
—	—	Survival product	0.000	0.000

tion error is repaired, but discrimination is not: mean pooled AUC at $k=5$ stays near 0.61, because the calibrators create ties or reverse rankings on the shifted target. Arm B is source-dependent: with NASA-trained XGBoost and $k=5$, Severson AUC rises from 0.620 to 0.854, overshooting the within-LCO reference ceiling (ratio 1.04, an artifact of the mismatched ceiling rather than genuine super-ceiling recovery), while CALCE-trained updates leave the source operating point essentially untouched (LCO-holdout retention 0.997 to 0.996 for CALCE-trained XGBoost) and ALL-LCO updates damage it substantially (see Table 15). Tables 14 and 15 summarize both arms.

III-J. Statistical Evidence for the SOH Ablation

Table 16 states the central comparison with the cell as the experimental unit: the paired cell-level bootstrap interval on $\Delta\text{AUC} = \text{AUC}(\text{with SOH}) - \text{AUC}(\text{without SOH})$ for the ALL-LCO source, with cycle-level DeLong

Table 14: Arm A Recalibration on Severson Without SOH ($H=20$): Values Average the Three LCO Sources and Available Tree Models; ECE Uses 10 Equal-Width Probability Bins. From the Released Recalibration Study Artifacts.

k	Zero ECE	Iso. ECE	Platt ECE	Temp. ECE	Iso. AUC	Platt AUC	Temp. AUC
5	0.817	0.044	0.045	0.563	0.616	0.599	0.669
10	0.817	0.010	0.010	0.561	0.634	0.646	0.651
20	0.817	0.013	0.013	0.560	0.634	0.655	0.648
40	0.817	0.015	0.015	0.561	0.637	0.657	0.653

Table 15: Arm B Recovery at $k = 5$ on Severson Without SOH, $H = 20$. Recovery Ratio Uses the Within-LCO Ceiling; Retention Is the LCO-Holdout AUC Before and After the Update. *Ratio above 1.0 is a ceiling-mismatch artifact, not super-ceiling recovery (see text).

Training	Model	Zero AUC	Arm B AUC	Recovery	DeLong p	LCO retention
NASA	XGBoost	0.620	0.854	+1.04*	$< 10^{-3}$	1.000 $\rightarrow$ 1.000
NASA	LightGBM	0.493	0.838	+0.88	$< 10^{-3}$	1.000 $\rightarrow$ 0.997
NASA	Random Forest	0.724	0.751	+0.19	$< 10^{-3}$	0.959 $\rightarrow$ 0.314
CALCE	XGBoost	0.686	0.796	+0.53	$< 10^{-3}$	0.997 $\rightarrow$ 0.996
ALL LCO	XGBoost	0.760	0.828	+0.47	$< 10^{-3}$	0.690 $\rightarrow$ 0.567

p -values as secondary evidence. On Severson, ΔAUC reaches 0.15 [0.10, 0.23] (XGBoost, $H=20$) with the interval far from zero; on Oxford the point estimates agree but paired intervals are undefined, so Oxford contributes directional consistency only. DeLong values underflow throughout and are reported as $p < 10^{-16}$ without interpretation (Section II-J). The claim rests on large bootstrapped effects on the 141-cell Severson target at every horizon, and on the failure-definition and same-chemistry controls of Sections III-F and III-G.

III-K. Operational Cost Analysis

AUC treats all errors as equal; a BMS does not. We select the threshold on source out-of-fold scores with a 10% false-alarm budget and apply it unchanged to the target (Table 17). On Severson with SOH, the transferred XGBoost threshold misses 32% of failures at 4% false alarms. The same protocol fails for the other ensembles with SOH: LightGBM alerts on 65% of negatives and Random Forest misses 100% of failures — threshold transfer fails for two of three ensembles even where ranking transfers. Without SOH, the XGBoost miss rate more than doubles to 70% at 10% false alarms; calibrated score types do not differ, because under shift the ranking is the

Table 16: SOH Ablation Evidence (ALL-LCO Source): Pooled AUC With and Without SOH, Paired Cell-Level Bootstrap 95% CI on the Difference (Primary), and Cycle-Level DeLong p -Value (Secondary; Inflated by Within-Cell Correlation).

Target	Model	H	AUC w/ SOH	AUC w/o	Delta AUC [boot 95% CI]	DeLong p
Oxford	XGBoost	10	0.980	0.433	+0.547 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	XGBoost	20	0.917	0.429	+0.488 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	XGBoost	30	0.991	0.377	+0.614 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	XGBoost	50	0.932	0.580	+0.351 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	LightGBM	10	0.996	0.547	+0.449 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	LightGBM	20	0.971	0.332	+0.639 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	LightGBM	30	0.743	0.484	+0.259 (CI undefined: 5 cells)	$9.7e - 13$
Oxford	LightGBM	50	0.832	0.448	+0.384 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	Random Forest	10	0.978	0.626	+0.352 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	Random Forest	20	0.989	0.581	+0.409 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	Random Forest	30	0.971	0.462	+0.508 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Oxford	Random Forest	50	0.933	0.460	+0.473 (CI undefined: 5 cells)	$< 10^{-16}$ (underflow)
Severson	XGBoost	10	0.878	0.741	+0.137 [0.079, 0.249]	$< 10^{-16}$ (underflow)
Severson	XGBoost	20	0.896	0.750	+0.146 [0.097, 0.233]	$< 10^{-16}$ (underflow)
Severson	XGBoost	30	0.906	0.745	+0.162 [0.115, 0.237]	$< 10^{-16}$ (underflow)
Severson	XGBoost	50	0.917	0.736	+0.180 [0.134, 0.250]	$< 10^{-16}$ (underflow)
Severson	LightGBM	10	0.870	0.736	+0.134 [0.105, 0.200]	$< 10^{-16}$ (underflow)
Severson	LightGBM	20	0.882	0.718	+0.164 [0.127, 0.227]	$< 10^{-16}$ (underflow)
Severson	LightGBM	30	0.872	0.771	+0.102 [0.030, 0.230]	$< 10^{-16}$ (underflow)
Severson	LightGBM	50	0.899	0.744	+0.155 [0.112, 0.214]	$< 10^{-16}$ (underflow)
Severson	Random Forest	10	0.903	0.740	+0.163 [0.133, 0.225]	$< 10^{-16}$ (underflow)
Severson	Random Forest	20	0.889	0.746	+0.143 [0.106, 0.226]	$< 10^{-16}$ (underflow)
Severson	Random Forest	30	0.919	0.748	+0.171 [0.138, 0.232]	$< 10^{-16}$ (underflow)
Severson	Random Forest	50	0.921	0.746	+0.175 [0.114, 0.264]	$< 10^{-16}$ (underflow)

Table 17: Operational Evaluation at H=20 (XGBoost, ALL-LCO Source): Decision Threshold Selected on Source Out-of-Fold Scores With a 10% False-Alarm Budget, Then Applied Unchanged to the Target. Reported Are the Target Miss Rate (FNR) and Target False-Alarm Rate (FPR).

Target	Features	raw		Platt		isotonic		temperature	
		FNR	FPR	FNR	FPR	FNR	FPR	FNR	FPR
Oxford	with SOH	0%	93%	0%	93%	0%	93%	0%	93%
Oxford	no SOH	0%	92%	0%	92%	0%	92%	0%	92%
Severson	with SOH	32%	4%	32%	4%	31%	6%	32%	4%
Severson	no SOH	70%	10%	70%	10%	67%	12%	70%	10%

asset and the ranking collapsed. On Oxford, nearly every row is positive, so no discriminating operating point exists there at all. The decision-curve view (Fig. 5) agrees: with SOH the model buys net benefit across thresholds; without SOH the curve hugs zero.




Figure 5: Decision-curve analysis (net benefit vs. threshold) for ALL-LCO → Severson without SOH at $H=20$: no score type produces a usable operating point once SOH is removed.

III-L. SHAP: The Shortcut Is Visible in the Model

TreeSHAP [15] gives a direct view of the mechanism. Figs. 6–8 show the attributions for the three tree models in NASA-to-Oxford transfer at $H=20$ with SOH kept: SOH dominates by a wide margin, cycle index a distant second, and the voltage, current, temperature, and duration features contribute little. Figs. 9–11 repeat the same comparisons with SOH removed, and every remaining feature collapses to near-zero SHAP spread — matching, in the model’s internal accounting, the quantitative collapse of Sections III-C–III-E and the sensor-inversion of Section III-D.

III-M. Embedded Deployment Validation

Three independent validation stages confirm that the C tree engine reproduces the Python training libraries through the output transforms of (8) and (9): a Python walker over the extracted trees, the same walker cross-checked on the full 1 028-row set, and the C engine run on-device on an ESP32-S3 against the library’s `predict_proba()`. Fig. 13 plots every one of the 1 028 on-device predictions against the Python values for all three models; the points fall on the identity line and the residual histograms are centered at zero. The measured on-device maximum absolute error is 2.4×10^{-7} for XGBoost, 1.1×10^{-6} for LightGBM, and 3.0×10^{-8} for Random Forest, all far inside the 10^{-5} tolerance.




Figure 6: SHAP feature importance for XGBoost in NASA-to-Oxford transfer at $H=20$, with SOH. SOH dominates by a wide margin.

Table 18 lists sizes, validation error, and inference time; Fig. 12 plots latency and footprint. The ensembles occupy about 372 kB. At 240 MHz: from flash, 0.76, 1.55, and 2.36 ms per model; from PSRAM, 0.50, 0.93, and 1.29 ms; single-model from SRAM, 0.32, 0.43, and 0.50 ms. The full 372 kB binary does not fit internal SRAM (free heap roughly 320 kB), but each model does, and the firmware runs one model at a time. Timings land above the 150–250 μ s projection because the working set exceeds the ESP32-S3’s data cache; even so, the all-three PSRAM path clears the 100 ms per-cycle budget by roughly $37\times$. The deployment result is deliberately independent of the validity question: whatever the models are reading, on this dataset they can read it in under a millisecond on a \$12 microcontroller. The on-device outputs are raw ensemble probabilities through (8)–(9) with no calibration layer: the Platt and isotonic maps live off-device, and Section III-I shows source-fitted maps do not survive dataset shift, so a deployed system must




Figure 7: SHAP feature importance for LightGBM in NASA-to-Oxford transfer at H=20, with SOH. SOH dominates by a wide margin.

Table 18: Embedded Deployment: Model Size, Accuracy, On-Device Validation Error, and Measured Inference Time. Errors Are Maximum $|C \text{ engine} - \text{predict_proba()}|$ Over All 1 028 Rows. Timings Measured On-Device at 240 MHz With the Binary in Flash (PROGMEM), PSRAM, or a Single Model in Internal SRAM.

Model	Nodes	Max Err	Flash	PSRAM	SRAM (1 model)
XGBoost	3 272	2.4×10^{-7}	0.76 ms	0.50 ms	0.32 ms
LightGBM	7 000	1.1×10^{-6}	1.55 ms	0.93 ms	0.43 ms
Random Forest	16 064	3.0×10^{-8}	2.36 ms	1.29 ms	0.50 ms
All three	26 336	—	4.7 ms	2.7 ms	—

revalidate or refit its calibrator on target-labeled cells rather than trust the laboratory calibration.




Figure 8: SHAP feature importance for Random Forest in NASA-to-Oxford transfer at $H=20$, with SOH. SOH dominates by a wide margin.

III-N. Replication on New Chemistries: NCA, NMC, and LFP Transfer Targets

The transfer claim is re-tested on 60 cells of NMC, NCA and LFP chemistry from two further laboratories (Battery Archive groups 5–8 of Section II-B), with the LCO training pool and protocol untouched — the new cells never enter model selection — and the experiment run on the same protocol version as every result above. Fig. 14 shows the post-cleaning capacity traces with the pre-registered audit decisions. Table 19 reports the NASA+CALCE $\rightarrow$ new-target transfer at $H=20$, with and without SOH, together with the paired cell-level bootstrap for the SOH ablation.

With SOH available, transfer from the LCO sources spans pooled AUC 0.77–1.00 across the new targets — the same upper-bound regime as Table 4. Removing SOH drops transfer to 0.36–1.00, and the paired per-cell bootstrap puts the with-minus-without difference at 0.00–0.64 AUC across the three en-




Figure 9: SHAP feature importance for XGBoost in NASA-to-Oxford transfer at $H=20$, without SOH. All features collapse to near-zero spread, so no transferable signal remains.

sembles — on the SNL groups the LCO-trained ensemble contributes nothing portable on NCA, NMC, or LFP chemistry from an unseen laboratory, exactly as it contributed nothing on the earlier LFP targets. The HNEI NMC group is the exception: full no-SOH transfer (cycle index included) stays at 0.99 pooled there. All 15 HNEI cells cycle under one condition and one fixed charge protocol, so cycle-index and voltage cues survive the laboratory change; this is the dataset-shift mechanism of Section III-G operating in reverse — narrow, stereotyped protocols transfer, diverse ones do not — not chemistry-specific knowledge. The common-set column repeats the check on the unit-safe feature subset (Section II-C), controlling for the structurally missing current and temperature columns of the export: the SNL collapse is not an artifact of the zero-imputed columns — though only for the available telemetry, per the Limitations. On the new groups the voltage-sag endpoint never fires before the SOH endpoint (0 of 60 admitted cells), so the composite




Figure 10: SHAP feature importance for LightGBM in NASA-to-Oxford transfer at $H=20$, without SOH. All features collapse to near-zero spread.

label is, effectively, the SOH criterion alone; the failure-definition ablation of Section III-F therefore cannot be re-run on these targets, and all with-SOH numbers here are to be read as label-proximity upper bounds, as defined in Section II-A.

III-O. Operating-Condition Shift at Fixed Chemistry and Laboratory

The same-chemistry controls of Section III-G varied laboratory while holding chemistry fixed; here the complement is measured — hold chemistry *and* laboratory fixed and vary only the operating condition. For each SNL group (NCA, NMC, LFP), the test cell is chosen from an operating condition (temperature band, C-rate, or both) that contributes no training cell in that fold; metadata parsed from cell identifiers provides the condition labels and is never used as a feature. Table 20 compares this leave-condition-out protocol with the ordinary mixed-cell folds on the same group.




Figure 11: SHAP feature importance for Random Forest in NASA-to-Oxford transfer at H=20, without SOH. All features collapse to near-zero spread.

Holding out an entire operating condition at fixed chemistry and laboratory costs little with SOH and more without it: on SNL-NCA the with-SOH pooled AUC moves only from 0.99 in ordinary mixed-cell folds to 0.99 under the joint holdout, and across all interpretable group-axis combinations the largest with-SOH drop is 0.02 AUC points — whereas the sensor-only model degrades before the with-SOH one does, falling to 0.89 on NCA (both axes) and 0.67 on NMC. Two caveats bound the claim’s strength. First, the mixed-fold raw reference on the 6-cell SNL-NMC group is itself unstable — with six cells the split composition dominates, and raw AUC there (0.65 with SOH) disagrees with the 0.87 that the same model reaches under the paper’s cell-disjoint outer folds — so its reference column is marked — and no drop is claimed for that group. Second, and most importantly, the distance-rule baseline is again undisturbed: its pooled AUC under condition holdout spans 0.886–0.999, essentially its mixed-fold value, so whatever the




Figure 12: Deployment trade-offs: (a) measured single-row inference time per model at 240 MHz with the binary in memory-mapped flash, PSRAM, or a single model in internal SRAM (dashed line: all three models from PSRAM); (b) per-model and combined binary footprint as a share of the ESP32-S3’s roughly 320 kB of free internal SRAM — each model fits, the combined 372 kB binary does not.

Table 19: Transfer to New Chemistries at H=20 (NASA+CALCE Source). Pooled AUC With Cell-Level Bootstrap 95% CI; Δ Is the Paired Per-Cell Bootstrap Difference (With – Without SOH). New Groups Are Transfer Targets Only.

Target	Model	With SOH	Without SOH	Delta (paired)
HNEI NMC	XGBoost	0.993 [0.983,0.999]	0.989 [0.977,0.997]	+0.004
HNEI NMC	LightGBM	0.999 [0.998,0.999]	0.988 [0.974,0.996]	+0.011
HNEI NMC	Random Forest	0.999 [0.999,1.000]	0.996 [0.994,0.999]	+0.003
SNL NCA	XGBoost	0.972 [0.956,0.988]	0.878 [0.824,0.930]	+0.094
SNL NCA	LightGBM	0.985 [0.974,0.993]	0.850 [0.791,0.911]	+0.135
SNL NCA	Random Forest	0.983 [0.974,0.992]	0.911 [0.875,0.951]	+0.072
SNL NMC	XGBoost	0.772 [0.677,0.979]	0.709 [0.612,0.953]	+0.063
SNL NMC	LightGBM	0.847 [0.743,0.982]	0.693 [0.570,0.959]	+0.153
SNL NMC	Random Forest	0.858 [0.748,0.975]	0.742 [0.689,0.874]	+0.115
SNL LFP	XGBoost	0.998 [0.997,1.000]	0.496 [0.458,0.531]	+0.502
SNL LFP	LightGBM	0.999 [0.997,1.000]	0.359 [0.279,0.443]	+0.640
SNL LFP	Random Forest	0.999 [0.998,1.000]	0.666 [0.569,0.769]	+0.332

rule keys on is invariant to operating condition. Two conclusions follow. The residual within-laboratory signal that survives dataset shift is only mildly condition-dependent — chemistry and laboratory shift, not protocol condition, carry the larger penalty — and no competing shortcut of comparable size hides in temperature or C-rate: with-SOH scores remain well above their no-SOH counterparts under condition holdout, so the SOH coordinate, not




Figure 13: On-device C engine predictions versus Python `predict_proba()` over all 1 028 rows and three models (agreement scatter on the identity line, inset residual histogram of C minus Python in units of 10^{-6}). Maximum absolute errors are 2.4×10^{-7} (XGBoost), 1.1×10^{-6} (LightGBM), and 3.0×10^{-8} (Random Forest).

Table 20: Leave-Condition-Out Versus Mixed-Cell Folds at H=20 (XGBoost). Conditions Are Temperature Band, C-Rate, or Both; Metadata Never Enters the Features. Ref. Is the Ordinary Mixed-Cell AUC for the Same Feature Set (— Where Unstable or Not Run).

Group	Held-out axis	With SOH	Ref.	No SOH	Ref.
SNL NCA	temp $\times$ rate	0.989	0.985	0.894	0.928
SNL NCA	temperature	0.989	0.985	0.888	0.928
SNL NCA	C-rate	0.961	0.985	0.820	0.928
SNL NMC	temp $\times$ rate	0.847	—	0.666	—
SNL NMC	temperature	0.787	—	0.569	—
SNL NMC	C-rate	0.847	—	0.666	—
SNL LFP	temp $\times$ rate	0.991	0.994	—	—
SNL LFP	temperature	0.995	0.994	—	—
SNL LFP	C-rate	0.994	0.994	—	—

an operating-condition coordinate, is the portable quantity — portable for the definitional reason quantified in Section III-F. The HNEI group is omitted because its cells cycle under a single condition (25 °C, mixed fixed discharge C-rates under one charge protocol), leaving no condition axis to hold out. SNL LFP was run with the with-SOH and common no-SOH feature sets only, so its full-feature no-SOH column reads — by design, not by omission.

III-P. Key Findings

- Within-distribution models are reliable (fold-mean AUC 0.90–0.96, Platt-calibrated) and run on-device in under a millisecond per SRAM-resident model.
- Cross-dataset transfer collapses without SOH (Severson up to 0.90 with SOH to 0.53–0.75 without); an unfitted SOH distance rule matches the ensembles.
- With-SOH numbers are diagnostic-proximity upper bounds: prognosis-only evaluation gives 0.838 [0.760, 0.991] on Severson and 0.73–0.99 on the new targets.
- Dataset and laboratory shift govern the collapse at least as much as chemistry; narrow protocols transfer, diverse ones do not, with no temperature or C-rate shortcut in the available telemetry.
- Source-fitted calibration and transferred thresholds do not survive shift; target recalibration repairs ECE but not ranking.

IV. Discussion, Analysis, and Novelty

IV-A. The SOH Shortcut, Precisely Stated

The evidence assembled above supports a precise statement. When a failure-risk model trained on one battery dataset is evaluated on another, its apparent transfer performance is dominated by the SOH coordinate; much of that dominance is definitional, because SOH enters the failure label itself, and the failure-definition ablation of Section III-F measures how much: relabeling failure by voltage sag alone removes about a fifth of the with-SOH transfer advantage on Severson. What remains is a weak, real health-state signal; everything else the sensors measure is dataset-specific, and the factorial ablation shows no non-SOH feature group carries usable cross-dataset signal.

The replication of Section III-N extends each of these statements to data the pipeline never had access to when it was designed. “SOH is a chemistry-specific lookup key,” the earlier framing, claimed more than the design could identify; the defensible claim is that apparent cross-dataset transferability is largely a health-state shortcut whose size this paper measures with label-, feature-, and dataset-level controls.

IV-B. Implications for Battery Management Practitioners

Four practical points follow:

- **1)** Within-chemistry, within-laboratory models from standard charge-discharge features are reliable (fold-mean AUC above 0.89, Platt-calibrated) and run on a \$12 microcontroller in under a millisecond per SRAM-resident model. Deploy them where the training distribution holds.
- **2)** Do not transfer models across datasets on the strength of an AUC table. If SOH is available at the target, an unfitted threshold rule is the correct first baseline; if SOH is unavailable, expect chance-level performance unless the target protocol is as narrow as the source.
- **3)** Quote risks from a coherent formulation: four independent classifiers can rank “risk within 10 cycles” above “risk within 50 cycles,” while the survival product of Section II-G cannot.
- **4)** Under shift, prefer raw or Platt-scaled scores, expect calibration to need labeled target data, and choose thresholds on source data with an explicit cost model — then verify.

IV-C. Limitations

The honesty of this paper has limits of its own. Oxford contributes five cells — one with a single positive row — at horizon-invariant 100-cycle logging resolution, so it is a pooled-detection target re-verified on Severson. CALCE contributes seven cells, so its leave-one-cell-out uncertainty is wide. The Battery Archive groups sharpen the chemistry claim but not the condition claim: SNL NCA/NMC span few conditions with as few as three cells each, and mostly-censored SNL LFP is supporting evidence. The export lacks per-cycle current and temperature; affected features are zero-imputed, conclusions rest on the common feature set, and the sensor collapse cannot

rule out richer-telemetry signals. The small three-seed GRU convicts no architecture beyond this instance, was not re-evaluated on the Battery Archive targets, and might differ from pretrained, transformer, or state-space models. The hazard model scores post-failure rows like every formulation. Cost analysis fixes 20:1 with source-selected thresholds. Transfer remains one-directional (LCO sources); isolating chemistry needs same-laboratory cross-chemistry campaigns.

V. Conclusion

This study widened a multi-horizon failure-risk framework and found the widening demanded a validity analysis. Within dataset the framework is sound: fold-mean AUC 0.90–0.96 under cell-disjoint evaluation, a monotone discrete-time hazard formulation, and an about-372 kB deployment reproducing the training libraries to 1.1×10^{-6} on a \$12 ESP32-S3.

Across datasets, the apparent transferability is mostly a State-of-Health shortcut. A distance-to-threshold rule on SOH alone reproduces the tuned ensembles’ transfer performance; removing SOH collapses every model class toward or below chance; relabeling failure by voltage sag alone — a criterion SOH does not define — removes about a fifth of the with-SOH advantage; and same-chemistry cross-laboratory controls show the collapse is caused by dataset shift at least as much as by chemistry. The supported conclusion is deliberately narrower than “SOH is chemistry-specific”: much of the observed cross-dataset discrimination is attributable to SOH, partly because SOH defines the failure criterion itself. Crucially, these statements replicate on 60 NMC, NCA and LFP cells from two further laboratories: with-SOH transfer spans 0.77–1.00 and drops to 0.36–0.91 without SOH outside the single-condition HNEI group (paired deltas 0.06–0.64) — a pre-registered replication, not a post-hoc fit.

Calibration behaves as the validity story predicts: Platt preserves discrimination best at comparable calibration error, isotonic fails under shift, labeled-target recalibration repairs calibration error but not discrimination, and no score type yields a usable operating point without SOH. Evidence is carried by cell-level bootstrap intervals, with cycle-level tests demoted to a consistency check.

Future work, in order of directness: an SOH-free failure definition as the primary cross-dataset target; same-laboratory, cross-chemistry aging campaigns; learned dataset-invariant representations; cell-level hierarchical cali-

bration; and multi-seed, multi-architecture sequence-model studies at generalizable scale.




Figure 14: Data audit of the Battery Archive groups after the pre-registered cleaning rule (formation-cycle drop, Hampel despiking, level-shift truncation, voltage validity, minimum 20 usable cycles): (a) admitted-cell capacity traces with the 80% end-of-life criterion; (b) retained versus excluded cycles per group; (c) recorded exclusion reasons. Of 75 downloaded cells, 60 are admitted and 15 are excluded for exactly one recorded reason.

Acknowledgment

This work received no external funding. The authors declare no conflicts of interest. The source code, data, trained models, firmware, and validation pipeline are publicly available at <https://github.com/touhidsiddiqueeraj-bit/Multi-Horizon-Hazard-Models-for-Battery-Failure-Prediction>.

References

- [1] H. Meng and Y. Li, “A review on prognostics and health management (PHM) methods of lithium-ion batteries,” *Renew. Sustain. Energy Rev.*, vol. 116, Art. no. 109405, 2019.
- [2] D. Roman, S. Saxena, V. Robu, et al., “Machine learning pipeline for battery state-of-health estimation,” *Nat. Mach. Intell.*, vol. 3, no. 4, pp. 447–456, 2021.
- [3] V. Sulzer, P. Mohtat, A. Aitio, et al., “The challenge and opportunity of battery lifetime prediction from field data,” *Joule*, vol. 5, no. 8, pp. 1934–1955, 2021.
- [4] L. Breiman, “Random forests,” *Mach. Learn.*, vol. 45, no. 1, pp. 5–32, 2001.
- [5] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. ACM SIGKDD*, 2016, pp. 785–794.
- [6] G. Ke, Q. Meng, T. Finley, et al., “LightGBM: A highly efficient gradient boosting decision tree,” in *Proc. NeurIPS*, vol. 30, 2017, pp. 3146–3154.
- [7] T. A. Shikdar and H. Laaksonen, “Learning when not to use a battery: Multihorizon failure intelligence,” *Int. Trans. Electr. Energy Syst.*, vol. 36, Art. no. 6000810, 2026.
- [8] S. Sahoo, K. S. Hariharan, S. Agarwal, et al., “Transfer learning based generalized framework for state of health estimation of Li-ion cells,” *Sci. Rep.*, vol. 12, Art. no. 13173, 2022.
- [9] J. Lu, R. Xiong, J. Tian, C. Wang, and F. Sun, “Deep learning to estimate lithium-ion battery state of health without additional degradation experiments,” *Nat. Commun.*, vol. 14, Art. no. 2760, 2023.

- [10] J. C. Platt, “Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods,” in *Adv. Large Margin Classifiers*, MIT Press, 1999, pp. 61–74.
- [11] B. Zadrozny and C. Elkan, “Transforming classifier scores into accurate multiclass probability estimates,” in *Proc. ACM SIGKDD*, 2002, pp. 694–699.
- [12] A. Niculescu-Mizil and R. Caruana, “Predicting good probabilities with supervised learning,” in *Proc. ICML*, 2005, pp. 625–632.
- [13] L. Huang, J. Zhao, B. Zhu, H. Chen, and S. K. L. M. v. Broucke, “An experimental investigation of calibration techniques for imbalanced data,” *IEEE Access*, vol. 8, pp. 127343–127352, 2020.
- [14] T. A. Shikdar and H. Laaksonen, “How far does the model hold? Robustness of multihorizon battery failure intelligence under realistic operating conditions,” submitted for publication, 2026.
- [15] S. M. Lundberg, G. G. Erion, H. Chen, et al., “From local explanations to global understanding with explainable AI for trees,” *Nat. Mach. Intell.*, vol. 2, no. 1, pp. 56–67, 2020.
- [16] P. Grzesik and D. Mrozek, “Combining machine learning and edge computing: Opportunities, challenges, platforms, frameworks, and use cases,” *Electronics*, vol. 13, no. 3, Art. no. 640, 2024.
- [17] C. Gupta and A. Ramdas, “Online Platt scaling with calibrating,” in *Proc. ICML*, PMLR 202, 2023, pp. 12182–12204.
- [18] B. Saha and K. Goebel, “Battery data set,” NASA Ames Prognostics Data Repository, 2007.
- [19] CALCE Battery Research Group, “Battery aging datasets,” Univ. Maryland, 2023.
- [20] C. R. Birkel and D. A. Howey, “Oxford battery degradation dataset 1,” Univ. Oxford, 2017.
- [21] K. A. Severson, P. M. Attia, N. Jin, et al., “Data-driven prediction of battery cycle life before capacity degradation,” *Nature Energy*, vol. 4, pp. 383–391, 2019.

- [22] P. M. Attia, A. Grover, N. Jin, et al., “Closed-loop optimization of fast-charging protocols for batteries with machine learning,” *Nature*, vol. 578, no. 7795, pp. 397–402, 2020.
- [23] Battery Archive, “Open battery degradation testing data,” 2023. [Online]. Available: <https://batteryarchive.org>. Consolidated cell cycling records from the Hawaii Natural Energy Institute and Sandia National Laboratories.
- [24] Y. Preger, H. C. Barkholtz, A. Fresquez, et al., “Degradation of commercial lithium-ion cells as a function of chemistry and cycling conditions,” *J. Electrochem. Soc.*, vol. 167, no. 12, Art. no. 120532, 2020.
- [25] E. R. DeLong, D. M. DeLong, and D. L. Clarke-Pearson, “Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach,” *Biometrics*, vol. 44, no. 3, pp. 837–845, 1988.